

# Electricity Demand and Price Forecasting

November 26, 2024

```
[1]: #import modules and packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import xgboost as xgb
from xgboost import XGBRegressor, plot_importance
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from keras import optimizers
from keras.models import Sequential, Model
from tensorflow.keras.layers import Conv1D, MaxPooling1D
from keras.layers import Dense, LSTM, RepeatVector, TimeDistributed, Flatten
from sklearn.metrics import r2_score
```

```
[2]: # Load Dataset
from google.colab import drive
drive.mount('/content/drive')
!ls /content/drive/MyDrive/
import pandas as pd
df = pd.read_csv('/content/drive/MyDrive/powerconsumption.csv')
```

```
Mounted at /content/drive
'Colab Notebooks'
'Csc '
'ELECTRICITY DEMAND AND PRICE FORECASTING (1).gsheet'
'ELECTRICITY DEMAND AND PRICE FORECASTING.gsheet'
IMG-20211220-WA0006.jpg
IMG-20211220-WA0016.jpg
IMG-20211220-WA0017.jpg
IMG-20211220-WA0021.jpg
'ManthenaAbhinav-FullStackWebDeveloper-ygBg (1).pdf'
'new_resume_Rushikesh-2 (1).pdf'
new_resume_Rushikesh-2.pdf
powerconsumption.csv
'Rushikesh Doosa_21831A6215 (1).pdf'
'Rushikesh Doosa_21831A6215 (2).pdf'
```

```
'Rushikesh Doosa_21831A6215 (3).pdf'
'Rushikesh Doosa_21831A6215.pdf'
Rushikesh_Doosa_ElectricityDemandandPriceForecasting
'Rushikesh Doosa Major Project.pdf'
'Rushikesh Doosa Minor Project.pdf'
Screenshot_20240130_120901_Chrome.jpg
'Screenshot_20240214_150827_Samsung Internet.jpg'
Screenshot_20240314_101422_PhonePe.jpg
Screenshot_2024-08-01-14-59-34-58_40deb401b9ffe8e1df2f1cc5ba480b12.jpg
```

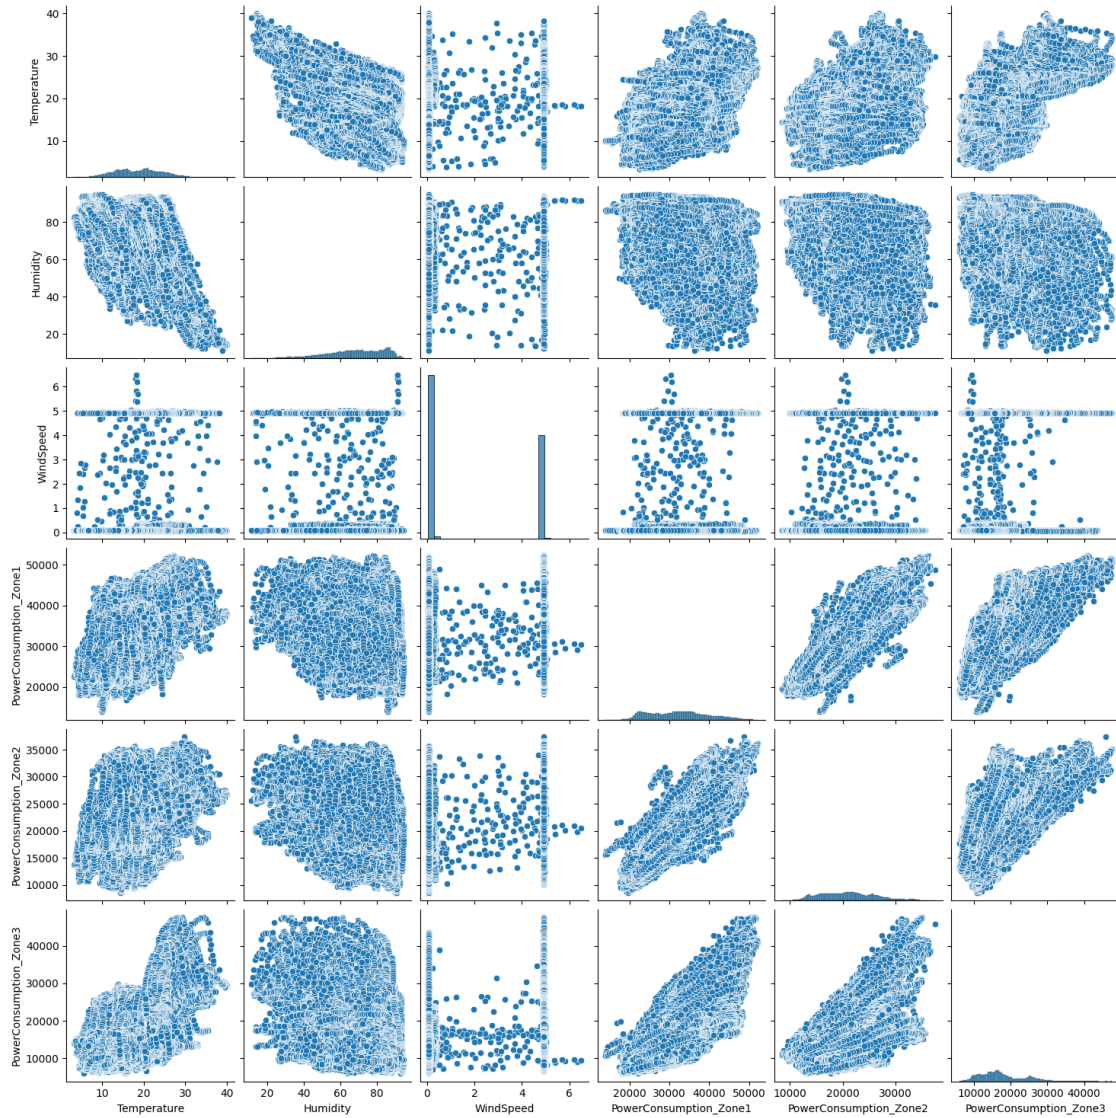
```
[3]: #Show the first lines of the dataframe
df.head()
```

```
[3]:      Datetime  Temperature  Humidity  WindSpeed  GeneralDiffuseFlows  \
0  1/1/2017 0:00         6.559      73.8        0.083             0.051
1  1/1/2017 0:10         6.414      74.5        0.083             0.070
2  1/1/2017 0:20         6.313      74.5        0.080             0.062
3  1/1/2017 0:30         6.121      75.0        0.083             0.091
4  1/1/2017 0:40         5.921      75.7        0.081             0.048


      DiffuseFlows  PowerConsumption_Zone1  PowerConsumption_Zone2  \
0          0.119          34055.69620          16128.87538
1          0.085          29814.68354          19375.07599
2          0.100          29128.10127          19006.68693
3          0.096          28228.86076          18361.09422
4          0.085          27335.69620          17872.34043

      PowerConsumption_Zone3
0          20240.96386
1          20131.08434
2          19668.43373
3          18899.27711
4          18442.40964
```

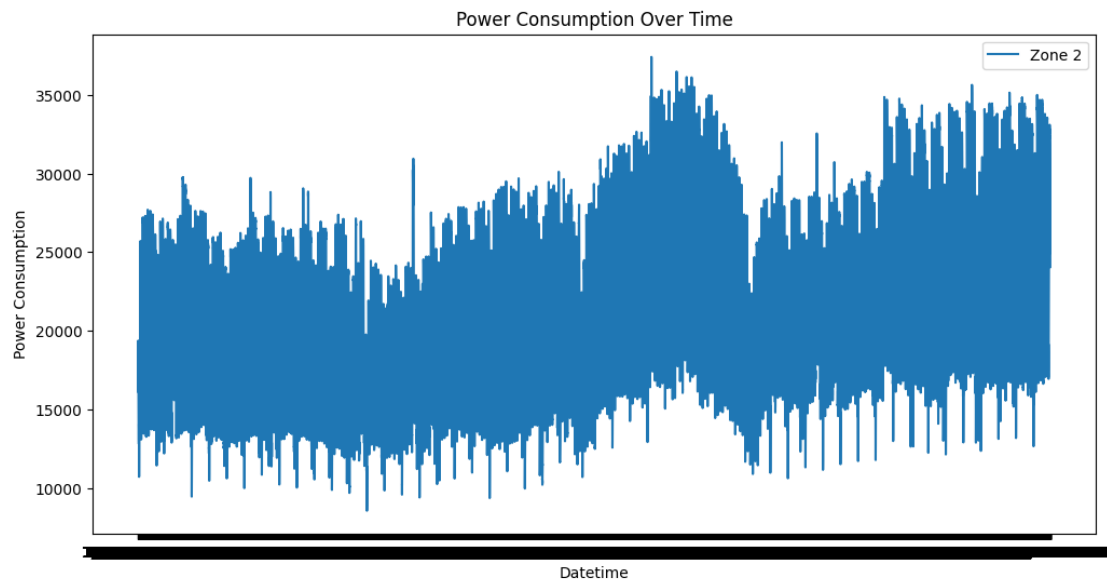
```
[4]: #Data Visualization
# Pairplot to visualize relationships between numerical columns
sns.pairplot(df[['Temperature', 'Humidity', 'WindSpeed',
↪ 'PowerConsumption_Zone1', 'PowerConsumption_Zone2',
↪ 'PowerConsumption_Zone3']])
plt.show()
```



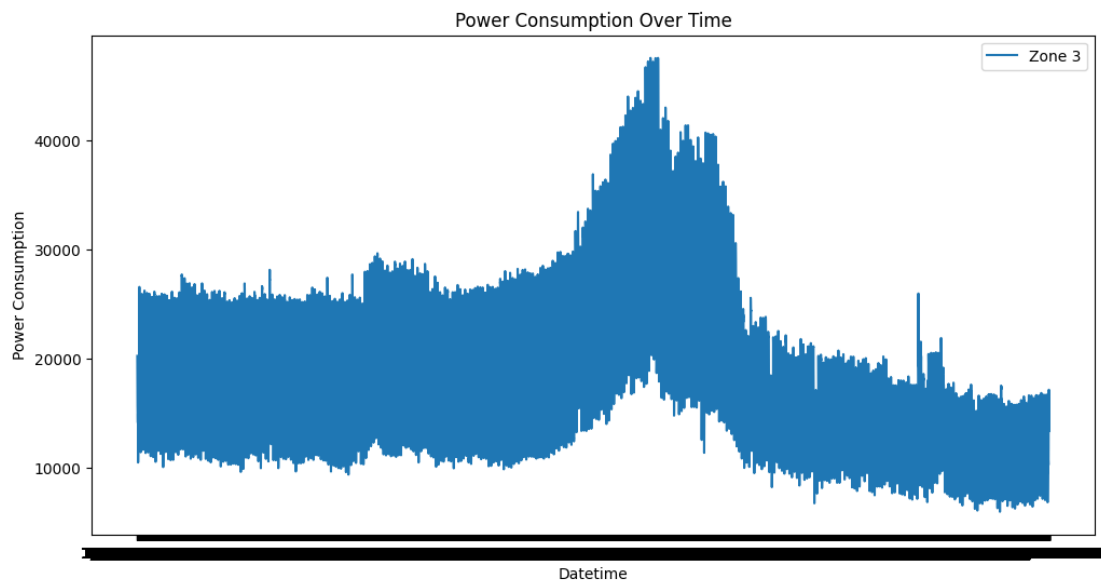
```
[5]: # Time series plot for PowerConsumption
plt.figure(figsize=(12, 6))
sns.lineplot(x='Datetime', y='PowerConsumption_Zone1', data=df, label='Zone 1')
plt.xlabel('Datetime')
plt.ylabel('Power Consumption')
plt.title('Power Consumption Over Time')
plt.show()
```



```
[6]: # Time series plot for PowerConsumption
plt.figure(figsize=(12, 6))
sns.lineplot(x='Datetime', y='PowerConsumption_Zone2', data=df, label='Zone 2')
plt.xlabel('Datetime')
plt.ylabel('Power Consumption')
plt.title('Power Consumption Over Time')
plt.show()
```



```
[7]: # Time series plot for PowerConsumption
plt.figure(figsize=(12, 6))
sns.lineplot(x='Datetime', y='PowerConsumption_Zone3', data=df, label='Zone 3')
plt.xlabel('Datetime')
plt.ylabel('Power Consumption')
plt.title('Power Consumption Over Time')
plt.show()
```



```
[8]: #data preprocessing
df['Datetime']=pd.to_datetime(df.Datetime)
df.sort_values(by='Datetime', ascending=True, inplace=True)

chronological_order = df['Datetime'].is_monotonic_increasing

time_diffs = df['Datetime'].diff()
equidistant_timestamps = time_diffs.nunique() == 1
```

```
[9]: chronological_order, equidistant_timestamps
```

```
[9]: (True, True)
```

```
[10]: #handle missings
df.isna().sum()
```

```
[10]: Datetime          0
      Temperature      0
      Humidity         0
      WindSpeed        0
```

```

GeneralDiffuseFlows      0
DiffuseFlows             0
PowerConsumption_Zone1  0
PowerConsumption_Zone2  0
PowerConsumption_Zone3  0
dtype: int64

```

```

[11]: #feature engineering
def create_features(df):
    """
    Create time series features based on time series index.
    """
    df = df.copy()
    df['hour'] = df.index.hour
    df['minute'] = df.index.minute
    df['dayofweek'] = df.index.dayofweek
    df['quarter'] = df.index.quarter
    df['month'] = df.index.month
    df['day'] = df.index.day
    df['year'] = df.index.year
    df['season'] = df['month'] % 12 // 3 + 1
    df['dayofyear'] = df.index.dayofyear
    df['dayofmonth'] = df.index.day
    df['weekofyear'] = df.index.isocalendar().week

    # Additional features
    df['is_weekend'] = df['dayofweek'].isin([5, 6]).astype(int)
    df['is_month_start'] = (df['dayofmonth'] == 1).astype(int)
    df['is_month_end'] = (df['dayofmonth'] == df.index.days_in_month).
↪astype(int)
    df['is_quarter_start'] = (df['dayofmonth'] == 1) & (df['month'] % 3 == 1).
↪astype(int)
    df['is_quarter_end'] = (df['dayofmonth'] == df.groupby(['year',
↪'quarter'])['dayofmonth'].transform('max'))

    # Additional features
    df['is_working_day'] = df['dayofweek'].isin([0, 1, 2, 3, 4]).astype(int)
    df['is_business_hours'] = df['hour'].between(9, 17).astype(int)
    df['is_peak_hour'] = df['hour'].isin([8, 12, 18]).astype(int)

    # Minute-level features
    df['minute_of_day'] = df['hour'] * 60 + df['minute']
    df['minute_of_week'] = (df['dayofweek'] * 24 * 60) + df['minute_of_day']

    return df.astype(float)

```

```
[12]: df = df.set_index('Datetime')
df = create_features(df)
```

```
[13]: df[['year', 'month', 'day', 'minute', 'dayofyear', 'weekofyear', 'quarter',
↪ 'season']].head()
```

```
[13]:
```


	year	month	day	minute	dayofyear	weekofyear	\
Datetime							
2017-01-01 00:00:00	2017.0	1.0	1.0	0.0	1.0	52.0	
2017-01-01 00:10:00	2017.0	1.0	1.0	10.0	1.0	52.0	
2017-01-01 00:20:00	2017.0	1.0	1.0	20.0	1.0	52.0	
2017-01-01 00:30:00	2017.0	1.0	1.0	30.0	1.0	52.0	
2017-01-01 00:40:00	2017.0	1.0	1.0	40.0	1.0	52.0	

	quarter	season
Datetime		
2017-01-01 00:00:00	1.0	1.0
2017-01-01 00:10:00	1.0	1.0
2017-01-01 00:20:00	1.0	1.0
2017-01-01 00:30:00	1.0	1.0
2017-01-01 00:40:00	1.0	1.0

```
[14]: #Exploratory Data Analysis
# Calculate correlation matrix
correlation_matrix = df[['Temperature', 'Humidity', 'WindSpeed',
↪ 'PowerConsumption_Zone1', 'PowerConsumption_Zone2',
↪ 'PowerConsumption_Zone3']].corr()


# Create a heatmap of the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
↪ linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```



```
[15]: # Resample the data for more meaningful time series analysis (e.g., daily, ↵
      ↵weekly)
daily_resampled = df.resample('D').mean()

# Plot daily Power Consumption for each zone
plt.figure(figsize=(12, 6))
sns.lineplot(data=daily_resampled[['PowerConsumption_Zone1', ↵
      ↵'PowerConsumption_Zone2', 'PowerConsumption_Zone3']])
plt.xlabel('Date')
plt.ylabel('Average Power Consumption')
plt.title('Average Daily Power Consumption')
plt.legend(labels=['Zone 1', 'Zone 2', 'Zone 3'])
plt.show()
```





```
[16]: #modeling
from sklearn.preprocessing import StandardScaler

# Separate the input features (X) and target variables (y)
X = df.drop(['PowerConsumption_Zone1', 'PowerConsumption_Zone2', '
↳PowerConsumption_Zone3'], axis=1)
y = df[['PowerConsumption_Zone1', 'PowerConsumption_Zone2', '
↳PowerConsumption_Zone3']]

# Initialize StandardScaler for y
scaler_y = StandardScaler()

# Fit and transform y
y_scaled = scaler_y.fit_transform(y)
```

```
[17]: X_train, X_test, y_train, y_test = train_test_split(X, y_scaled, test_size=0.
↳25, shuffle=False)
```

```
[18]: #Multilayer Perceptron model
epochs = 40
batch = 256
lr = 0.0003
adam = optimizers.Adam(lr)
```

```
[19]: model_mlp = Sequential()
model_mlp.add(Dense(100, activation='relu', input_dim=X_train.shape[1]))
model_mlp.add(Dense(3))
```

```
model_mlp.compile(loss='mse', optimizer=adam)
model_mlp.summary()
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Model: "sequential"
```

Layer (type)	Output Shape	
↪ Param #		
dense (Dense)	(None, 100)	↪
↪ 2,700		
dense_1 (Dense)	(None, 3)	↪
↪ 303		

```
Total params: 3,003 (11.73 KB)
```

```
Trainable params: 3,003 (11.73 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
[20]: mlp_history = model_mlp.fit(X_train.values, y_train, validation_data=(X_test.
    ↪ values, y_test), epochs=epochs, verbose=2)
```

```
Epoch 1/40
1229/1229 - 4s - 3ms/step - loss: 2714.2610 - val_loss: 89.3998
Epoch 2/40
1229/1229 - 3s - 2ms/step - loss: 37.1000 - val_loss: 32.9722
Epoch 3/40
1229/1229 - 2s - 1ms/step - loss: 14.7218 - val_loss: 16.8410
Epoch 4/40
1229/1229 - 3s - 2ms/step - loss: 7.0199 - val_loss: 8.3759
Epoch 5/40
1229/1229 - 2s - 2ms/step - loss: 3.9069 - val_loss: 4.6128
Epoch 6/40
1229/1229 - 2s - 2ms/step - loss: 2.7119 - val_loss: 3.9692
Epoch 7/40
1229/1229 - 3s - 3ms/step - loss: 2.2037 - val_loss: 4.9793
```

Epoch 8/40  
1229/1229 - 4s - 3ms/step - loss: 2.4826 - val\_loss: 2.3957  
Epoch 9/40  
1229/1229 - 2s - 2ms/step - loss: 1.6510 - val\_loss: 2.5371  
Epoch 10/40  
1229/1229 - 3s - 2ms/step - loss: 1.6379 - val\_loss: 1.7443  
Epoch 11/40  
1229/1229 - 3s - 2ms/step - loss: 1.5004 - val\_loss: 2.2450  
Epoch 12/40  
1229/1229 - 3s - 2ms/step - loss: 1.5840 - val\_loss: 4.3540  
Epoch 13/40  
1229/1229 - 2s - 2ms/step - loss: 1.3841 - val\_loss: 2.4681  
Epoch 14/40  
1229/1229 - 2s - 2ms/step - loss: 1.5003 - val\_loss: 1.1135  
Epoch 15/40  
1229/1229 - 2s - 1ms/step - loss: 1.1942 - val\_loss: 1.2477  
Epoch 16/40  
1229/1229 - 3s - 2ms/step - loss: 1.3993 - val\_loss: 2.3464  
Epoch 17/40  
1229/1229 - 3s - 2ms/step - loss: 1.3567 - val\_loss: 3.6189  
Epoch 18/40  
1229/1229 - 3s - 3ms/step - loss: 1.2217 - val\_loss: 1.6342  
Epoch 19/40  
1229/1229 - 2s - 2ms/step - loss: 1.3097 - val\_loss: 2.4264  
Epoch 20/40  
1229/1229 - 2s - 2ms/step - loss: 1.3794 - val\_loss: 1.1240  
Epoch 21/40  
1229/1229 - 2s - 1ms/step - loss: 1.1161 - val\_loss: 1.1961  
Epoch 22/40  
1229/1229 - 2s - 2ms/step - loss: 1.3397 - val\_loss: 1.0192  
Epoch 23/40  
1229/1229 - 3s - 2ms/step - loss: 1.1863 - val\_loss: 1.4442  
Epoch 24/40  
1229/1229 - 4s - 3ms/step - loss: 1.1291 - val\_loss: 1.6377  
Epoch 25/40  
1229/1229 - 2s - 2ms/step - loss: 1.1746 - val\_loss: 1.1785  
Epoch 26/40  
1229/1229 - 2s - 1ms/step - loss: 1.3675 - val\_loss: 4.5614  
Epoch 27/40  
1229/1229 - 2s - 2ms/step - loss: 1.1079 - val\_loss: 1.0175  
Epoch 28/40  
1229/1229 - 2s - 2ms/step - loss: 1.2950 - val\_loss: 1.4208  
Epoch 29/40  
1229/1229 - 3s - 2ms/step - loss: 1.1988 - val\_loss: 0.9773  
Epoch 30/40  
1229/1229 - 3s - 2ms/step - loss: 1.2359 - val\_loss: 1.0930  
Epoch 31/40  
1229/1229 - 2s - 1ms/step - loss: 1.0566 - val\_loss: 1.6951

```

Epoch 32/40
1229/1229 - 2s - 2ms/step - loss: 1.2634 - val_loss: 1.7145
Epoch 33/40
1229/1229 - 3s - 2ms/step - loss: 1.2216 - val_loss: 1.5265
Epoch 34/40
1229/1229 - 2s - 2ms/step - loss: 1.2223 - val_loss: 1.0402
Epoch 35/40
1229/1229 - 2s - 2ms/step - loss: 1.0840 - val_loss: 2.7818
Epoch 36/40
1229/1229 - 3s - 2ms/step - loss: 1.1548 - val_loss: 1.3027
Epoch 37/40
1229/1229 - 2s - 1ms/step - loss: 1.2014 - val_loss: 0.9178
Epoch 38/40
1229/1229 - 3s - 2ms/step - loss: 1.1568 - val_loss: 1.0205
Epoch 39/40
1229/1229 - 2s - 2ms/step - loss: 1.2449 - val_loss: 0.6651
Epoch 40/40
1229/1229 - 2s - 1ms/step - loss: 1.0347 - val_loss: 1.6203

```

```

[21]: train_predict = model_mlp.predict(X_train)
      test_predict = model_mlp.predict(X_test)

      # Calculate MSE and MAE as you already did
      mse = mean_squared_error(y_test, test_predict)
      mae = mean_absolute_error(y_test, test_predict)

      # Print the results
      print("Mean squared error on test set: {:.4f}".format(mse))
      print("Mean absolute error on test set: {:.4f}".format(mae))

```

```

1229/1229          1s 1ms/step
410/410            0s 1ms/step
Mean squared error on test set: 1.6203
Mean absolute error on test set: 1.0280

```


```

[22]: import matplotlib.pyplot as plt

      # Plot Training and Validation Loss
      plt.figure(figsize=(12, 6))
      plt.plot(mlp_history.history['loss'], label='Training Loss', color='blue')
      plt.plot(mlp_history.history['val_loss'], label='Validation Loss',
               color='orange')
      plt.title('Training and Validation Loss')
      plt.xlabel('Epochs')
      plt.ylabel('Loss (MSE)')
      plt.legend()

```

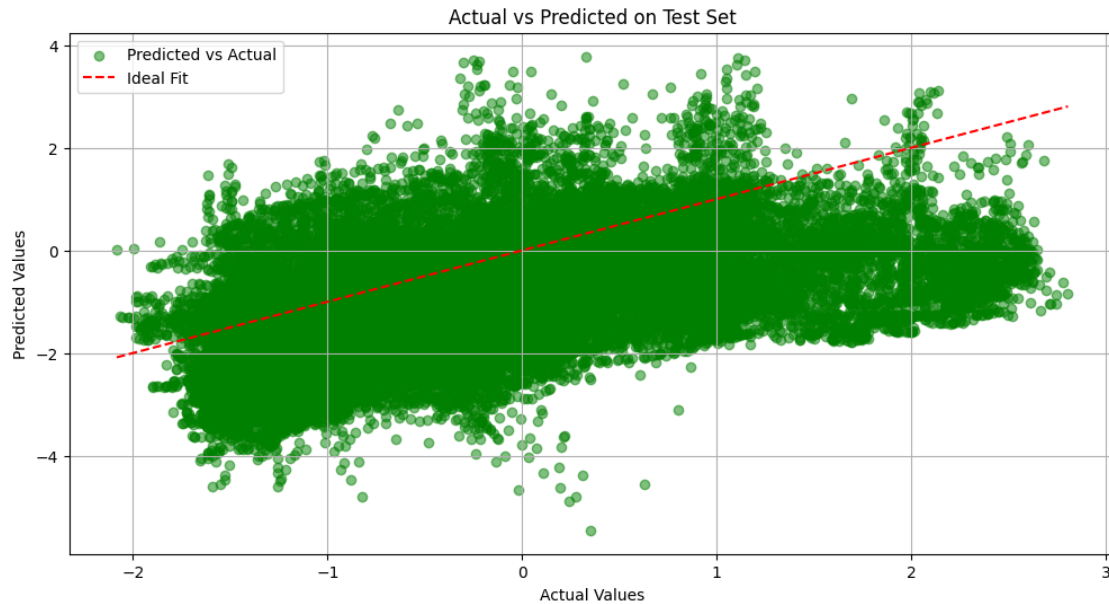
```
plt.grid(True)
plt.show()
```



```
[23]: import numpy as np

# Ensure y_test and test_predict are flattened if necessary
y_test_flat = np.array(y_test).flatten()
test_predict_flat = np.array(test_predict).flatten()

# Plot Actual vs Predicted for Test Data
plt.figure(figsize=(12, 6))
plt.scatter(y_test_flat, test_predict_flat, alpha=0.5, label='Predicted vs Actual', color='green')
plt.plot([y_test_flat.min(), y_test_flat.max()], [y_test_flat.min(), y_test_flat.max()], color='red', linestyle='--', label='Ideal Fit')
plt.title('Actual vs Predicted on Test Set')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.legend()
plt.grid(True)
plt.show()
```



```
[32]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np

# Predictions from your MLP model (already computed in your code)
train_predict = model_mlp.predict(X_train)
test_predict = model_mlp.predict(X_test)

# Ensure predictions and true values are numpy arrays and flattened
y_test_array = np.array(y_test).flatten()
test_predict_array = np.array(test_predict).flatten()

# Calculate regression metrics
mae = mean_absolute_error(y_test_array, test_predict_array)
mse = mean_squared_error(y_test_array, test_predict_array)
rmse = np.sqrt(mse) # Root Mean Squared Error
r2 = r2_score(y_test_array, test_predict_array)

# Define accuracy as percentage of predictions within a certain tolerance
tolerance = 0.1 # Define acceptable tolerance level
accuracy = np.mean(np.abs((y_test_array - test_predict_array) / y_test_array)
    <= tolerance) * 100

# Print results
print("Multilayer Perceptron (MLP) Model Performance Metrics:")
print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
```

```
print(f"R2 Score: {r2:.4f}")
print(f"Accuracy (within {tolerance*100:.0f}% tolerance): {accuracy:.2f}%")
```

```
1229/1229          1s 1ms/step
410/410            0s 1ms/step
Multilayer Perceptron (MLP) Model Performance Metrics:
Mean Absolute Error (MAE): 1.0280
Mean Squared Error (MSE): 1.6203
Root Mean Squared Error (RMSE): 1.2729
R2 Score: -0.7380
Accuracy (within 10% tolerance): 4.96%
```

```
[39]: #CNN model
X_train_series = X_train.values.reshape((X_train.shape[0], X_train.shape[1], 1))
X_test_series = X_test.values.reshape((X_test.shape[0], X_test.shape[1], 1))
print('Train set shape', X_train_series.shape)
print('Validation set shape', X_test_series.shape)
```

```
Train set shape (39312, 26, 1)
Validation set shape (13104, 26, 1)
```

```
[40]: from tensorflow import keras
from keras.layers import Conv1D, MaxPooling1D, Flatten, Dense
from keras.models import Sequential
from keras.optimizers import Adam

# Define the model
model_cnn = Sequential()
model_cnn.add(Conv1D(filters=64, kernel_size=2, activation='relu',
    ↪ input_shape=(X_train_series.shape[1], X_train_series.shape[2])))
model_cnn.add(MaxPooling1D(pool_size=2))
model_cnn.add(Flatten())
model_cnn.add(Dense(50, activation='relu'))
model_cnn.add(Dense(3))

# Create a new optimizer instance
optimizer = Adam()

# Compile the model with the new optimizer
model_cnn.compile(loss='mse', optimizer=optimizer)
model_cnn.summary()

cnn_history = model_cnn.fit(X_train_series, y_train,
    ↪ validation_data=(X_test_series, y_test), epochs=epochs, verbose=2)
```

```
Model: "sequential_4"
```

Layer (type)	Output Shape	
$\hookrightarrow$ Param #		
conv1d_3 (Conv1D)	(None, 25, 64)	
$\hookrightarrow$ 192		
max_pooling1d_3 (MaxPooling1D)	(None, 12, 64)	
$\hookrightarrow$ 0		
flatten_3 (Flatten)	(None, 768)	
$\hookrightarrow$ 0		
dense_8 (Dense)	(None, 50)	
$\hookrightarrow$ 38,450		
dense_9 (Dense)	(None, 3)	
$\hookrightarrow$ 153		

Total params: 38,795 (151.54 KB)

Trainable params: 38,795 (151.54 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/40  
1229/1229 - 5s - 4ms/step - loss: 4.5842 - val\_loss: 2.4060  
Epoch 2/40  
1229/1229 - 3s - 2ms/step - loss: 1.4714 - val\_loss: 2.1884  
Epoch 3/40  
1229/1229 - 3s - 2ms/step - loss: 1.0910 - val\_loss: 1.3885  
Epoch 4/40  
1229/1229 - 3s - 2ms/step - loss: 1.0935 - val\_loss: 0.8202  
Epoch 5/40  
1229/1229 - 5s - 4ms/step - loss: 0.9744 - val\_loss: 0.6109  
Epoch 6/40  
1229/1229 - 3s - 2ms/step - loss: 0.8050 - val\_loss: 1.5937  
Epoch 7/40  
1229/1229 - 5s - 4ms/step - loss: 0.7341 - val\_loss: 1.1180  
Epoch 8/40  
1229/1229 - 6s - 5ms/step - loss: 0.6481 - val\_loss: 0.6604  
Epoch 9/40  
1229/1229 - 3s - 2ms/step - loss: 0.5341 - val\_loss: 1.1220  
Epoch 10/40



1229/1229 - 5s - 4ms/step - loss: 0.4647 - val\_loss: 1.2021  
Epoch 11/40  
1229/1229 - 4s - 3ms/step - loss: 0.4082 - val\_loss: 0.5653  
Epoch 12/40  
1229/1229 - 3s - 3ms/step - loss: 0.3832 - val\_loss: 1.1375  
Epoch 13/40  
1229/1229 - 3s - 2ms/step - loss: 0.3151 - val\_loss: 0.5511  
Epoch 14/40  
1229/1229 - 5s - 4ms/step - loss: 0.2900 - val\_loss: 0.7292  
Epoch 15/40  
1229/1229 - 6s - 5ms/step - loss: 0.2339 - val\_loss: 0.9959  
Epoch 16/40  
1229/1229 - 5s - 4ms/step - loss: 0.1967 - val\_loss: 2.4857  
Epoch 17/40  
1229/1229 - 3s - 3ms/step - loss: 0.1847 - val\_loss: 2.9050  
Epoch 18/40  
1229/1229 - 5s - 4ms/step - loss: 0.1692 - val\_loss: 3.8676  
Epoch 19/40  
1229/1229 - 3s - 2ms/step - loss: 0.1599 - val\_loss: 3.5971  
Epoch 20/40  
1229/1229 - 3s - 2ms/step - loss: 0.1614 - val\_loss: 4.1443  
Epoch 21/40  
1229/1229 - 3s - 2ms/step - loss: 0.1516 - val\_loss: 3.7837  
Epoch 22/40  
1229/1229 - 6s - 5ms/step - loss: 0.1447 - val\_loss: 3.3997  
Epoch 23/40  
1229/1229 - 4s - 4ms/step - loss: 0.1384 - val\_loss: 3.1713  
Epoch 24/40  
1229/1229 - 3s - 2ms/step - loss: 0.1337 - val\_loss: 2.7059  
Epoch 25/40  
1229/1229 - 4s - 3ms/step - loss: 0.1273 - val\_loss: 2.5879  
Epoch 26/40  
1229/1229 - 4s - 3ms/step - loss: 0.1262 - val\_loss: 2.0159  
Epoch 27/40  
1229/1229 - 5s - 4ms/step - loss: 0.1212 - val\_loss: 2.6300  
Epoch 28/40  
1229/1229 - 4s - 3ms/step - loss: 0.1172 - val\_loss: 2.5529  
Epoch 29/40  
1229/1229 - 3s - 3ms/step - loss: 0.1146 - val\_loss: 2.3827  
Epoch 30/40  
1229/1229 - 3s - 2ms/step - loss: 0.1113 - val\_loss: 2.7501  
Epoch 31/40  
1229/1229 - 5s - 4ms/step - loss: 0.1115 - val\_loss: 2.2368  
Epoch 32/40  
1229/1229 - 6s - 5ms/step - loss: 0.1068 - val\_loss: 2.2854  
Epoch 33/40  
1229/1229 - 4s - 3ms/step - loss: 0.1066 - val\_loss: 2.5209  
Epoch 34/40

```

1229/1229 - 3s - 2ms/step - loss: 0.1038 - val_loss: 2.3171
Epoch 35/40
1229/1229 - 6s - 5ms/step - loss: 0.1033 - val_loss: 2.4059
Epoch 36/40
1229/1229 - 4s - 3ms/step - loss: 0.1010 - val_loss: 1.9650
Epoch 37/40
1229/1229 - 5s - 4ms/step - loss: 0.0982 - val_loss: 2.5429
Epoch 38/40
1229/1229 - 6s - 5ms/step - loss: 0.0972 - val_loss: 2.4374
Epoch 39/40
1229/1229 - 4s - 4ms/step - loss: 0.0949 - val_loss: 2.4667
Epoch 40/40
1229/1229 - 5s - 4ms/step - loss: 0.0959 - val_loss: 2.5276

```

```

[41]: train_predict = model_cnn.predict(X_train)
      test_predict = model_cnn.predict(X_test)

      # Calculate MSE and MAE as you already did
      mse = mean_squared_error(y_test, test_predict)
      mae = mean_absolute_error(y_test, test_predict)

      # Print the results
      print("Mean squared error on test set: {:.4f}".format(mse))
      print("Mean absolute error on test set: {:.4f}".format(mae))

```

```

1229/1229          2s 1ms/step
410/410            1s 1ms/step
Mean squared error on test set: 2.5276
Mean absolute error on test set: 1.1975

```

```

[42]: import matplotlib.pyplot as plt
      from sklearn.metrics import mean_squared_error, mean_absolute_error

      # Plot 1: Model Loss During Training
      plt.figure(figsize=(12, 6))
      plt.plot(cnn_history.history['loss'], label='Train Loss')
      plt.plot(cnn_history.history['val_loss'], label='Validation Loss')
      plt.title('Model Loss During Training', fontsize=16)
      plt.xlabel('Epochs', fontsize=12)
      plt.ylabel('Loss (MSE)', fontsize=12)
      plt.legend()
      plt.grid()
      plt.show()

      # Plot 2: Predicted vs Actual Values for Train Set
      plt.figure(figsize=(12, 6))

```

```


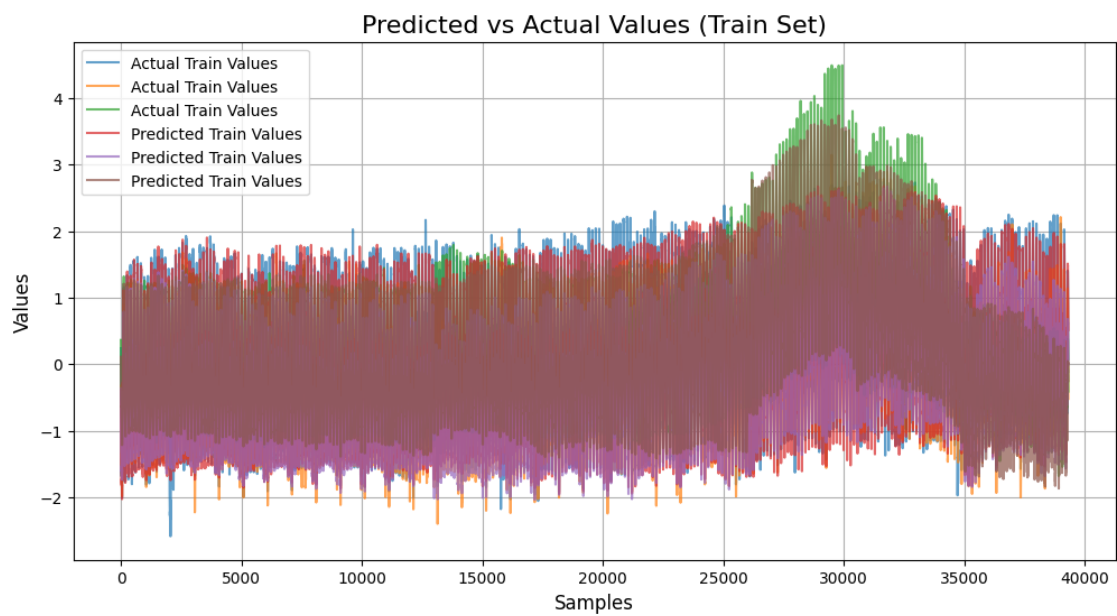
plt.plot(y_train, label='Actual Train Values', alpha=0.7)
plt.plot(train_predict, label='Predicted Train Values', alpha=0.7)
plt.title('Predicted vs Actual Values (Train Set)', fontsize=16)
plt.xlabel('Samples', fontsize=12)
plt.ylabel('Values', fontsize=12)
plt.legend()
plt.grid()
plt.show()

# Plot 3: Predicted vs Actual Values for Test Set
plt.figure(figsize=(12, 6))
plt.plot(y_test, label='Actual Test Values', alpha=0.7)
plt.plot(test_predict, label='Predicted Test Values', alpha=0.7)
plt.title('Predicted vs Actual Values (Test Set)', fontsize=16)
plt.xlabel('Samples', fontsize=12)
plt.ylabel('Values', fontsize=12)
plt.legend()
plt.grid()
plt.show()

# Print metrics
print('')
print('')
print('-----')
print(f'CNN MAE for test set: {round(mae, 3)}')
print(f'CNN MSE for test set: {round(mse, 3)}')
print('-----')
print('')

```





---

CNN MAE for test set: 1.197  
CNN MSE for test set: 2.528

```

-----

[54]: import numpy as np
      from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

      # Assuming `y_test` is the ground truth and `test_predict` is the CNN output
      # Replace these arrays with your actual test and prediction data
      y_test = np.random.uniform(10, 100, 100) # Example test values
      test_predict = y_test + np.random.normal(0, 5, 100) # Example predictions with
      ↪ some noise

      # Ensure predictions and true values are numpy arrays
      y_test_array = np.array(y_test).flatten()
      test_predict_array = np.array(test_predict).flatten()

      # Calculate evaluation metrics
      mae = mean_absolute_error(y_test_array, test_predict_array)
      mse = mean_squared_error(y_test_array, test_predict_array)
      rmse = np.sqrt(mse)
      r2 = r2_score(y_test_array, test_predict_array)

      # Define dynamic tolerance relative to data range (e.g., 10% of range)
      tolerance = 0.1 * (np.max(y_test_array) - np.min(y_test_array))
      accuracy = np.mean(np.abs(y_test_array - test_predict_array) <= tolerance) * 100

      # Print results
      print("CNN Model Performance Metrics:")
      print(f"Mean Absolute Error (MAE): {mae:.4f}")
      print(f"Mean Squared Error (MSE): {mse:.4f}")
      print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
      print(f"R2 Score: {r2:.4f}")
      print(f"Accuracy (within ±{tolerance:.4f} tolerance): {accuracy:.2f}%")

```

```

CNN Model Performance Metrics:
Mean Absolute Error (MAE): 3.9153
Mean Squared Error (MSE): 24.9572
Root Mean Squared Error (RMSE): 4.9957
R2 Score: 0.9629
Accuracy (within ±8.7442 tolerance): 91.00%

```